


# Programming distributed systems with NIO

---

Fabienne Boyer, LIG  
Université Joseph Fourier  
2014-2015



# Basic Socket IO




- Considering TCP
  - Lossless
  - Ordered
- Stream-oriented
  - Read one or more bytes at a time
- Blocking
  - *read()* blocks until there is some data to read
  - *write()* blocks until the data is written
  - **This generally requires one worker thread per socket (for reactivity)**
  - Not highly scalable

# New Socket IO (Java NIO)



- Buffer-oriented
  - Data is read from a stream into a buffer
  - Data is written from a buffer to a stream
- **Non-blocking**
  - *read()* only reads the data that is currently available
  - *write()* only writes the data that can be currently written
  - No need for several threads dedicated to blocking read/write operations on separate channels, a single thread (or a limited set of threads) can be used

# New Socket IO



- Typical use
  - IO intensive applications
  - Mono-threaded systems
  - Communication groups
- **Event-based programming model**
  - One (or more) thread waits for events (IO events and possibly others)
  - Any event is associated with a handler (also called *callback* or *reaction*)

# Event-based programming



## *Event-based*

Use a single thread  
(scheduler thread)

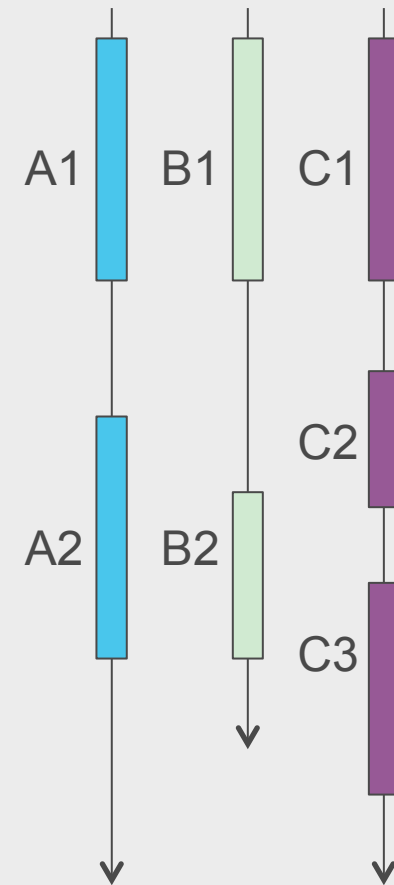
The code is  
programmed in small  
chunks of code that do  
not contain blocking  
instructions

No useless  
commutations

Fabienne Boyer, Basics of Distributed Programming



## *Thread-based*

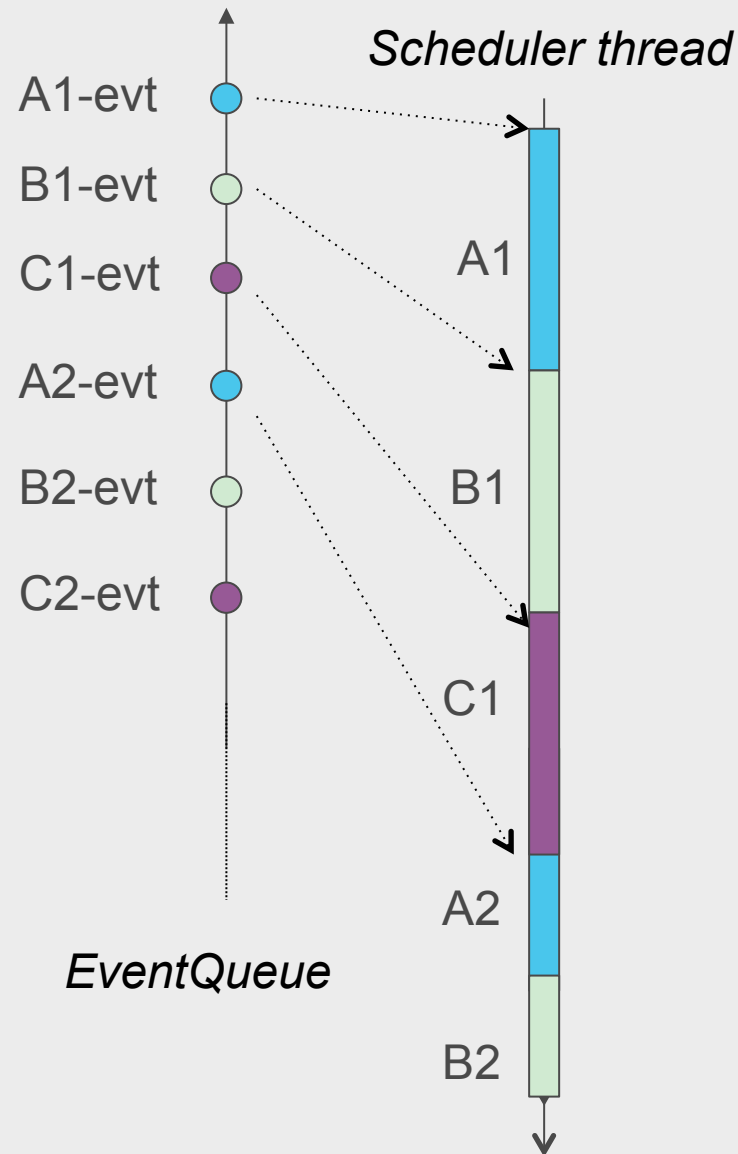


# Event-based programming



Events inform the scheduler of handlers that should be processed


The association between an event and an handler may be managed in several ways




# Event-based programming



Events can be fired by low-level signals (e.g., incoming data over an IO channel) or by handlers



# Event-based programming



## Focusing on task A

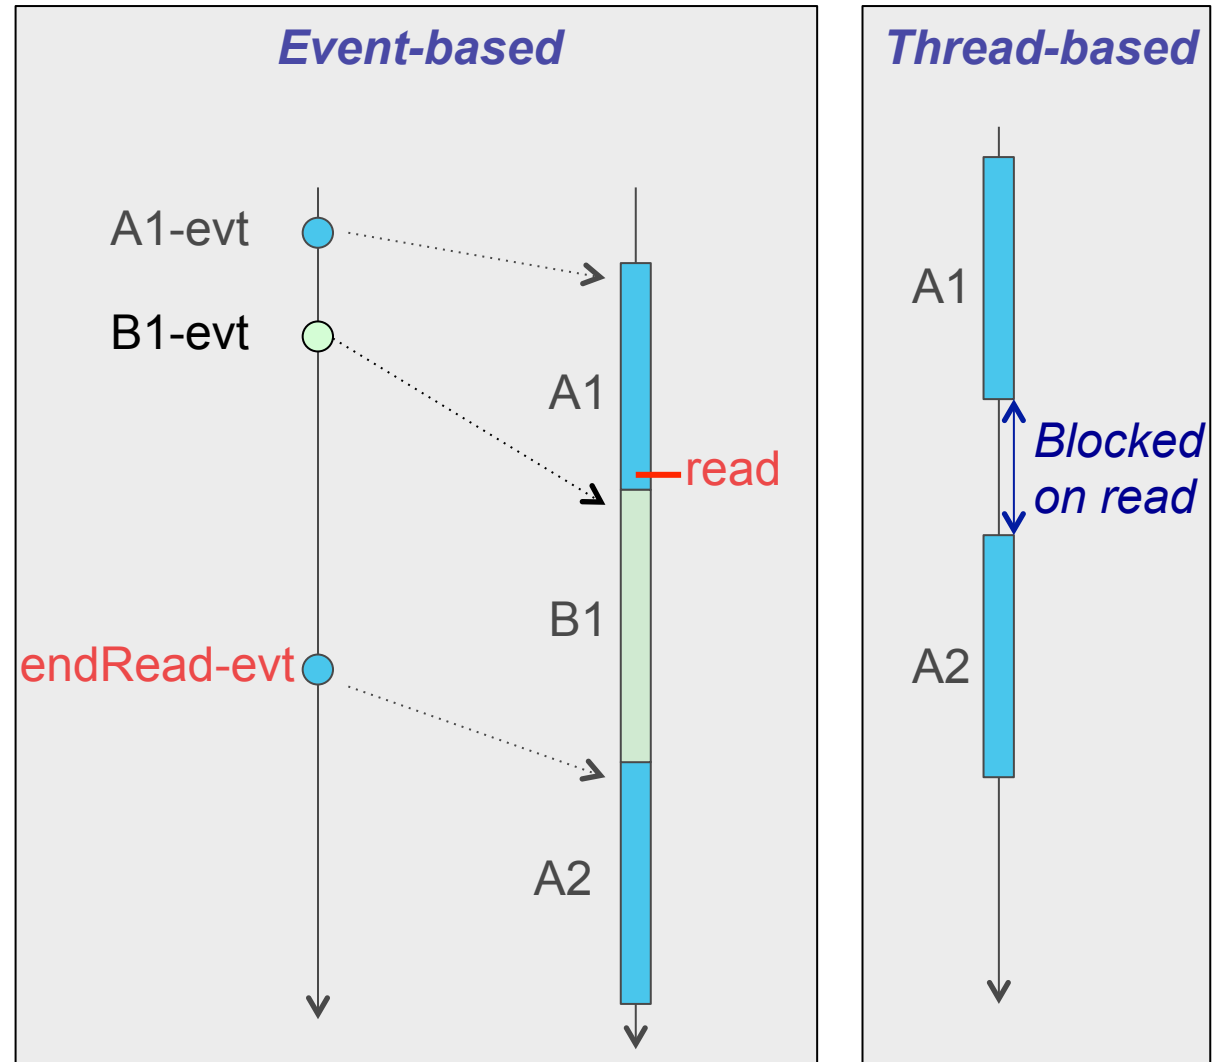
### (thread-based)

Suppose that the time between A1 and A2 is due to a blocking *read* operation

### (event-based)

We need a non-blocking read operation

Handling *endRead-evt* on task A leads to executing A2





# Event-based programming



## FIRST RULE:

**A handler runs to completion, it should not block**

*→ Note that blocking for a very short period of time may be allowed*

*→ Note that the scheduler may be multi-threaded for performance reasons but this is an implementation choice that should not change the programming model (except that reactions should be thread-safe)*

- So event programming does not work with classical sockets / streams
  - Sockets provide a blocking API (accept, connect, read, write)
  - Input and output streams are the same (read, write)

# Event-based programming



## SECOND RULE

**Handlers may have to be programmed as Automata (Finite State Machin)**

- When do we program a handler as an Automata?
  - When events do not always require the same treatment
  - The treatment depends on the previous events that were fired
  - Being given a current state and an event, the automata determines how the event should be processed
- We may use
  - A global automata for all events
  - An automata / event type (or per set of event types)
  - An automata / client
  - ..

# The GUI example



- GUI programming
  - Main (scheduler) thread goes in the Widget toolkit event pump
  - Only reacts to graphical events through callbacks
  - Callbacks are never called concurrently
  - Possible to submit events for later execution
  - Usually use an automata per widget

# NIO Programming Model



- NIO Events
  - ACCEPT, CONNECT, READ, WRITE
  - Fired by the NIO layer
- NIO Handlers
  - `handleAccept`, `handleConnect`, `handleRead`, `handleWrite`
- Use automata for
  - READ handler on a given channel, as we may need several read to compose a received message
  - WRITE handler on a given channel, as a single message may be sent in several times

# NIO Programming Model




- Programming handlers
  - Do not block within handlers
  - So we cannot program as follow

```
handleRead(byte[] b) {  
    int nbr = 0;  
    int len = b.length  
    while (nbr < len){  
        nbr += read(b, nbr, len-nbr)  
    }  
}
```

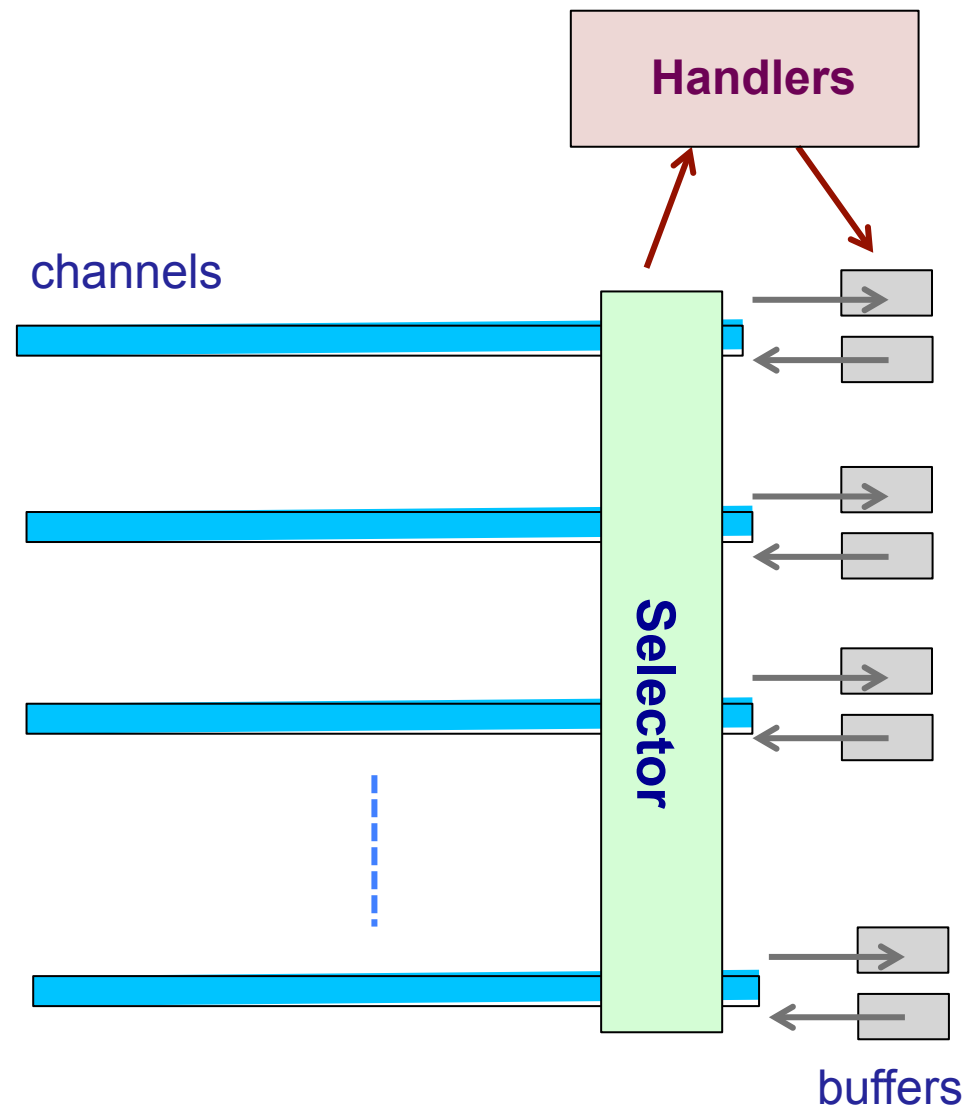
```
handleWrite(byte[] b) {  
    int nbw=0;  
    int len = b.length;  
    while (nbw < len){  
        nbw += write(b, nbw, len-nbw);  
    }  
}
```

- Basically
  - For read events on a given channel, use an automata that keeps *b* and *nbr*
  - For write events on a given channel, use an automata that keeps *b* and *nbw*

# NIO Concepts



- **Channels**  
Provide for data transfers between sockets and buffers
- **Buffers**  
Contiguous extent of memory for processing data
- **Selector**  
Allows to wait on several channels until one or more become available for data transfer



# NIO Channels



- To send or receive data
  - Lossless and FIFO
- Channels are close to streams except that
  - Channels always read to (or write from) a Buffer
  - `connect()`, `accept()`, `read()`, `write()` operations on Channels are asynchronous
- Some Channel implementations
  - *SocketChannel*, to read and write data over the network via TCP
  - *ServerSocketChannel*, to listen for TCP connections
    - For each incoming connection, a *SocketChannel* is created.

# NIO Buffers



- To hold data, either received or to be sent
  - Not a stream, a random-access buffer of a given size
  - Non-blocking send (send what can be sent)
  - Non blocking receive (only receive what is currently available)
  - Keep the memory of read / write indexes
- Provide optimized operations
  - Use the physical memory of the under-laying operating system for native IO operations
  - No additional copies when transfers are processed
- Buffer attributes
  - Capacity (number of bytes)
  - Position (next index for read/write)
  - Limit (maximum number of bytes that can be read / written)




# NIO Selector



- A selector allows a single thread to wait for IO events on multiple channels
  - `select()`, `select(long timeout)`, `selectNow()`
- Channels of interest should be registered on the selector
  - Register one or more channels
  - For each registered channel, get a selection key
  - Set interest on that key, any mix of `ACCEPT`, `CONNECT`, `READ`, `WRITE`
  - Interests on a key can be changed at any time
- Typical use
  - Always have `READ` to be notified
  - Only have `WRITE` when there is something to send
  - Only have `CONNECT` when you just connected a socket
  - Only have `ACCEPT` when you have created server sockets

# NIO Concepts



**REMIND TO  
NEVER BLOCK  
IN HANDLER !!**

## Handle accept

register the created channel on the selector

## Handle connect

register read/write interests


## Handle read

read bytes from channel into InBuffer and process received data

## Handle write

write bytes from OutBuffer into channel

# NIO – Server programming



```
public class NioServer implements Runnable {
    private InetAddress address;
    private int port, length = ...;
    private ServerSocketChannel serverChannel;
    private Selector selector;

    public NioServer(InetAddress hostAddress, int port) throws IOException {
        this.hostAddress = hostAddress;
        this.port = port;
        selector = SelectorProvider.provider().openSelector();

        // Create a new non-blocking server socket channel
        serverChannel = ServerSocketChannel.open();
        serverChannel.configureBlocking(false);
        serverChannel.socket().bind(new InetSocketAddress(hostAddress,port));

        // Be notified when connection requests arrive
        serverChannel.register(selector, SelectionKey.OP_ACCEPT);
    }

    public static void main(String[] args) {
        try { new Thread(new NioServer(null, 8888)).start(); } catch (IOException) ..
```

# NIO – Server programming




```
public void run() {
    while (true) {
        try {
            // Wait for an event one of the registered channels
            this.selector.select();

            // Some events have been received
            Iterator selectedKeys = this.selector.selectedKeys().iterator();
            while (selectedKeys.hasNext()) {
                SelectionKey key = (SelectionKey) selectedKeys.next();
                selectedKeys.remove();
                if (!key.isValid()) { continue; }

                // Handle the event
                if (key.isAcceptable()){           handleAccept(key);
                } else if (key.isConnectable()){ handleConnect(key);
                } else if (key.isReadable()){      handleRead(key);
                } else if (key.isWritable()) {      handleWrite(key);
            } } catch (Exception e) { ... } }
```

# NIO – Client programming



```
public class NioClient implements Runnable {
    private InetAddress address;
    private int port, length = ...;
    private SocketChannel clientChannel;
    private Selector selector;

    public NioClient(InetAddress hostAddress, int port) throws IOException {
        this.hostAddress = hostAddress;
        this.port = port;
        selector = SelectorProvider.provider().openSelector();
        // Create a new non-blocking socket channel
        clientChannel = SocketChannel.open();
        clientChannel.configureBlocking(false);
        // Be notified when connection is accepted
        clientChannel.register(selector, SelectionKey.OP_CONNECT);

        // Connect to the server
        clientChannel.connect(new InetSocketAddress(hostAddress, port));
    }

    public static void main(String[] args) {
        try { new Thread(new NioClient(null, 8888)).start(); } catch (IOException) ..
    }
}
```

# NIO – Client programming



```
public void run() {
    while (true) {
        try {
            // Wait for an event one of the registered channels
            this.selector.select();

            // Some events have been received
            Iterator selectedKeys = this.selector.selectedKeys().iterator();
            while (selectedKeys.hasNext()) {
                SelectionKey key = (SelectionKey) selectedKeys.next();
                selectedKeys.remove();
                if (!key.isValid()) { continue; }

                // Handle the event
                if (key.isAcceptable()){           handleAccept(key);
                } else if (key.isConnectable()){ handleConnect(key);
                } else if (key.isReadable()){      handleRead(key);
                } else if (key.isWritable()) {      handleWrite(key);
            } } catch (Exception e) { ... } }
```

# NIO - Connecting principles



- Connecting
  - Create a `SocketChannel`
  - Configure it to non-blocking
  - Configure it to TCP no delay
  - Register that channel on a selector, getting a key
  - Set the `CONNECT` interest
  - Start the connection operation (`SocketChannel.connect(InetSocketAddress)`)
- Later Connect Handling
  - When the key will be in the connectable state, the selector will unblock
  - The connection sequence must be finished (`SocketChannel.finishConnect()`)
  - Set the `READ` interest on the key to receive data

# NIO – Connect Handling



```
private void handleConnect(SelectionKey key) throws IOException {  
    SocketChannel socketChannel = (ServerSocketChannel) key.channel();  
  
    // finish the connection  
    socketChannel.finishConnect();  
  
    // register the read interest on the selector  
    socketChannel.register(this.selector, SelectionKey.OP_READ);  
  
    ..  
}
```



# NIO - Accepting principles



- Accepting
  - Create a ServerSocketChannel
  - Configure it to non-blocking
  - Bind it to a port
  - Register it to a selector, getting a key
  - Set at least the ACCEPT interest
- Later Accept Handling
  - Upon accepting a connection, the selector will unblock
  - The key will be in the acceptable state
  - Get the socket channel from the key
  - Set it to non blocking and noTCP delay
  - Set the READ interest on the key to receive data

# NIO – Accept Handling



```
private void handleAccept(SelectionKey key) throws IOException {  
  
    ServerSocketChannel serverSocketChannel = (ServerSocketChannel) key.channel();  
  
    // Accept the connection and make it non-blocking  
    SocketChannel socketChannel = serverSocketChannel.accept();  
    Socket socket = socketChannel.socket();  
    socketChannel.configureBlocking(false);  
  
    // Register the new SocketChannel with our Selector, indicating  
    // we'd like to be notified when there's data waiting to be read  
    socketChannel.register(this.selector, SelectionKey.OP_READ);  
}
```

# NIO – Reading (basic version)



- Reading
  - Set the READ interest on the channel
- Later Read Handling
  - Read available bytes and deliver them immediately

# NIO - Read Handling (basic version)



```
private void handleRead(SelectionKey key) throws IOException {
    SocketChannel socketChannel = (SocketChannel) key.channel();
    ByteBuffer inBuffer = ByteBuffer.allocate(length); // Read up to length bytes

    int nbread = 0;
    try {
        nbread = socketChannel.read(inBuffer);
    } catch (IOException e) {
        // the connection as been closed unexpectedly, cancel the selection and close the channel
        key.cancel();
        socketChannel.close();
        return;
    }
    if (nbread == -1) {
        // the socket has been shutdown remotely cleanly
        key.channel().close();
        key.cancel();
        return;
    }
    // process the received data, being aware that it may be incomplete
    deliver(this, socketChannel, inBuffer.array(), nbread);
}
```



# NIO – Writing (basic version)



- Writing
  - Manage a write buffer per channel
  - Put any message to send in this buffer (possibly losing part of the previous message if not entirely sent)
  - Set the WRITE interest on the channel
- Later writeHandler
  - Write as much as we can
  - Remove the WRITE interest if we wrote the entire buffer

# NIO – Writing (basic version)



```
// outBuffers contains the data to write per channel
Hashtable<SocketChannel, ByteBuffer> outBuffers = ...;

private void write(SocketChannel socketChannel, byte[] data) throws IOException {

    // we suppose that previous data in outBuffer have already been sent
    // or we do not mind loosing them
    outBuffers.put(socketChannel, ByteBuffer.wrap(data));

    // indicate we want to select OP-WRITE from now
    SelectionKey key = socketChannel.keyFor(this.selector);
    key.interestOps(SelectionKey.OP_READ | SelectionKey.OP_WRITE )
}
```



# NIO – Basic Write Handling



```
// outBuffers contains the data to write per channel
Hashtable<SocketChannel, ByteBuffer> outBuffers = ...;

private void handleWrite(SelectionKey key) throws IOException {
    SocketChannel socketChannel = (SocketChannel) key.channel();
    ByteBuffer outBuffer = outBuffers.get(socketChannel);
    if (outBuffer.remaining() > 0) {
        try {
            socketChannel.write(outBuffer);
        } catch (IOException e) {

            // The channel has been closed
            key.cancel();
            socketChannel.close();
            return;
        }
    } else ...
}
```

# Advanced Reading



- For any channel, we manage a read automata  
*Hashtable<SocketChannel, ReadAutomata> readAutomata = ...;*
- Read automata
  - Implements the *handleRead(..)* method to gather received bytes as they become available
  - Knows that each message is prefixed with its length on 4 bytes
  - Only deliver complete messages



# Read Automata



```
Class ReadAutomata {
    SocketChannel sock;
    ByteBuffer lenBuf = ByteBuffer.allocate(4); // for reading the length of a message
    ByteBuffer msgBuf = null; // for reading a message
    static final int READING_LENGTH = 1;
    static final int READING_MSG = 2;
    int currentState = READING_LENGTH; // initial state
    ...


    private void handleRead() throws IOException {
        if (state == READING_LENGTH){
            int nb = sock.read(lenBuf);
            ...
            if (lenBuf.remaining() == 0) {
                length = byteArrayToInt(lenBuf.array());
                msgBuf = ByteBuffer.allocate(byteArrayToInt(lengthBuf.array()));
                lenBuf.position(0);
                state = READING_MSG;
            }
        }
        if (state == READING_MSG) {
            ...
            if (msgBuf.remaining() == 0){ // the message has been fully received
                deliver(msgBuf.array()); // deliver it
                msgBuf = null;
            } else ...
        }
    }
}
```

# Advanced Writing




- For any channel, we manage a write automata  
*Hashtable<SocketChannel, WriteAutomata> writeAutomata =*
- Write automata
  - Manage a FIFO queue of messages to send
  - Prefix each message with its length when sending it
  - So each buffer to send contains the length of the message on 4 bytes

# Write Automata



```
Class WriteAutomata {
    SocketChannel sock;
    ArrayList<ByteBuffer> messages = new ArrayList<ByteBuffer>() ; // messages to send
    ByteBuffer lenBuf = ByteBuffer.allocate(4); // for writing the length of a message
    ByteBuffer msgBuf = null; // for writing a message
    static final int WRITING_LENGTH = 1;
    static final int WRITING_MSG = 2;
    int currentState = WRITING_LENGTH; // initial state
    ...
    private void handleWrite() throws IOException {
        if (state == WRITING_LENGTH){
            ...
            if (lenBuf.remaining() == 0) {
                state = WRITING_MSG;
            }
        } else if (state == WRITING_MSG) {
            if (msgBuf.remaining() > 0){
                nb = sock.write(msgBuf);
            }
            if (msgBuf.remaining() == 0){ // the message has been fully sent
                if (! messages.isEmpty()){
                    msgBuf = messages.remove(0);
                    lenBuf.position(0); lenBuf.putInt(0, buf.remaining());
                    state = WRITING_LENGTH;
                } else STATE = WRITING_DONE;
            }
        }
        ...
    }
}
```

# Write Automata



```
Class WriteAutomata {
    SocketChannel sock;
    ArrayList<ByteBuffer> messages = new ArrayList<ByteBuffer>() ; // messages to send
    ByteBuffer lenBuf = ByteBuffer.allocate(4); // for writing the length of a message
    ByteBuffer msgBuf = null; // for writing a message
    static final int WRITING_LENGTH = 1;
    static final int WRITING_MSG = 2;
    int currentState = WRITING_LENGTH; // initial state
    ...
    private void handleWrite() throws IOException {
        <see previous slide>
    }

    private void write(ByteBuffer msg) {
        messages.add(msg);
        if (state == WRITING_DONE)
            key.interestOps(SelectionKey.OP_READ | SelectionKey.OP_WRITE);
    }
}
```

# NIO – Work to do



- PingPong NIO server
  - A server waits for
    - Connection requests from clients
    - Messages sent by connected clients
    - The server returns to each client the message it sent
  - Once connected, a client spends its time sending textual messages and receiving the server's response
- First version: without automata management
- Second version: with automata management